


Sampling methods for counting temporal motifs


Paul Liu · Stanford University

Slides. tinyurl.com/wsdm19

Paper. [arXiv:1810.00980](https://arxiv.org/abs/1810.00980)



Joint work with
Austin Benson (Cornell) &
Moses Charikar (Stanford)



Stanford
University

Temporal network data is extremely common.



Private communication

e-mail, phone calls, text messages, instant messages




Payment systems

credit card transactions, cryptocurrencies, Venmo



Public communication

Q&A forums, Facebook walls, Wikipedia edits




Technical infrastructure

packets over the Internet, messages over supercomputer

Motifs, or small subgraph patterns, are commonly used to analyze static (non-temporal) networks.


1. Common feature for anomaly detection, role discovery, and other network machine learning problems.

[Noble-Cook 03; Sun+ 07; Henderson+ 12; Rohe-Qin 13; Rossi-Ahmed 15; Benson-Gleich-Leskovec 16]



2. Finding fundamental components of complex systems. [Milo+ 02]



- Triangles in social networks.
[Rapoport 53; Granovetter 73; Watts-Strogatz 98]
- Bi-directed length-2 paths in brain networks.
[Sporns-Kötter 04; Sporns+ 07; Honey+ 07]



Temporality adds important context to motifs.



Red edges occur before **blue** edges.

What is the difference between the following *temporal* motifs?



How does context change if $\delta = 1$ day? 1 month? 1 year?

[Paranjape-Benson-Leskovec 17] precisely defines the temporal motif counting problem.




Temporal network motif.

1. Directed multigraph with k edges
2. Edge ordering
3. Max. time span $\delta = 25$.

Motif instance.

k temporal edges that match the pattern that all occur within δ time.



There is a need to have scalable algorithms for real-time temporal motif analysis.

- Prior algorithms focused on enumeration [Mackey+ 18] or exact counts for small motifs [Paranjape-Benson-Leskovec 17].
- Algorithms were extremely memory extensive.
- Compute times on the order of days for our largest datasets, and could not be done in a streaming manner.

How do we enable real-time motif analysis for high-throughput temporal network data?

Majority of applications only
require *approximate* counts.


From exact to approximate: random sampling.

Idea. Use *importance sampling* to accelerate existing motif algorithms.

Our contributions.

- Theoretical foundation for temporal motif counting, showing that it is NP-hard.
- A sampling framework to accelerate existing temporal motif counting algorithms.
- New sampling algorithms are memory-efficient, and can be done in a streaming fashion.

Parallel sampling yields about two orders of magnitude speedup and enables otherwise infeasible computations.



Time scale $\delta = 1$ day.
16 threads.

dataset	# temporal edges	running time (seconds)			error
		exact	sampling	parallel sampling	
StackOverflow	47.9M	221.7	93.10	5.208	4.9%
EquinixChicago	345M	481.2	45.50	5.666	1.3%
RedditComments	636M	X	6739	2262	–

Using backtracking algorithm from [\[Mackey+ 18\]](#) as a sub-routine.

From exact to approximate: random sampling.

“Easy” solution: sample subset of edges from graph and run exact algorithm on it.

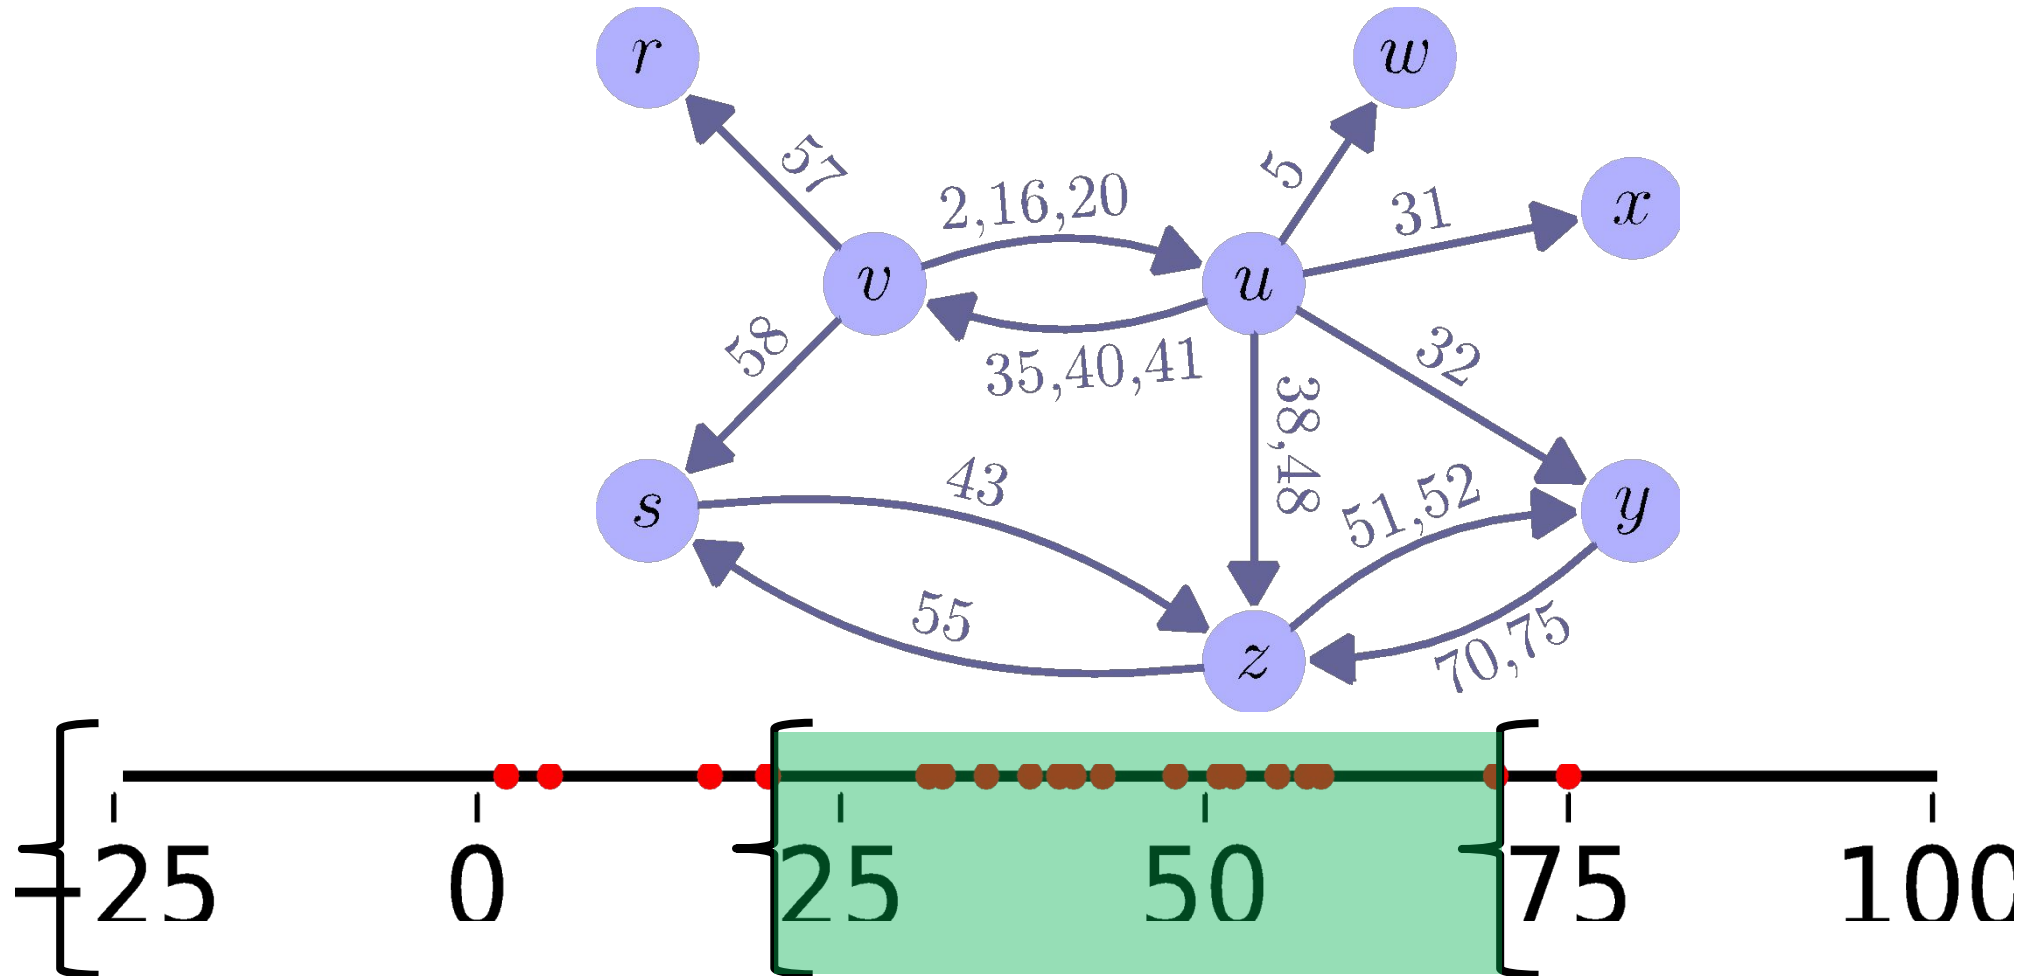
How do we sample? (Two simple ideas)

1. Sample uniformly random subset of edges? ✗ edges sampled often are too far apart in time.
2. Sample uniformly random windows of duration δ ? ✗ contribution from windows are uneven. May miss temporally dense regions of edges.

We find motifs in sampled windows and re-scale counts.

Sampling window length $w > \delta$ ($\delta = 25$, $w = 50$).


Choose random shift s uniformly from $[0, 1, \dots, w - 1]$ ($s = 20$).



We find motifs in sampled windows and re-scale counts.

Sampling window length $w > \delta$ ($\delta = 25, w = 50$).

Choose random shift s uniformly from $[0, 1, \dots, w - 1]$ ($s = 20$).




1. How do we re-scale exact counts?
2. Motifs can cross sampling intervals. How do we mitigate this?
3. How do we choose sampling probabilities q ?

We find motifs in sampled windows and re-scale counts.

Sampling window length $w > \delta$ ($\delta = 25$, $w = 50$).

Choose random shift s uniformly from $[0, 1, \dots, w - 1]$ ($s = 20$).



motif instance M


duration $d(M) = 32 - 16 = 16$

Theorem. If we sample window j with prob. q_j , then upscaling each found motif instance by $(1 - d(M) / w) / q_j$ is an unbiased estimator, where $d(M)$ is the duration of the motif instance M .

We find motifs in sampled windows and re-scale counts.

Sampling window length $w > \delta$ ($\delta = 25, w = 50$).

Choose random shift s uniformly from $[0, 1, \dots, w - 1]$ ($s = 20$).




- Using multiple random shifts and averaging the estimates reduces variance by capturing motifs that cross sampling intervals.
- $s = 20, s = 32, s = 37$
- Computation over each shift is parallelizable.

We find motifs in sampled windows and re-scale counts.

Sampling window length $w > \delta$ ($\delta = 25, w = 50$).


Choose random shift s uniformly from $[0, 1, \dots, w - 1]$ ($s = 20$).



- Set q_j for each shift. Larger $q_j \rightarrow$ more computation but less variance.
- **Importance sampling.** Only want to sample where motifs occur.
- **Heuristic.** Make q_j larger if more edges in sampling window.

Computation over windows is naturally *streaming*.

- Computations of different windows are independent.
- Old data from longer than one window ago can be thrown away.
- Multiple estimators can be run in parallel to ensure accuracy.



Our algorithm in a nutshell.

Input. Temporal motif and maximum time scale δ .

Output. Estimate of number of instances of the motif.

Partition data into windows

Choose a random shift s , and partition the data into windows aligned to s .

Parallel Importance Sampling

For j th window, sample with probability q_j .

Count and upscale

Upscale counts of motif instances depending on their duration $d(M)$.

Our algorithm in a nutshell.

Input. Temporal motif and maximum time scale δ .

Output. Estimate of number of instances of the motif.

Partition data into windows

Parallel Importance
Sampling

Count and upscale

Key advantages.


- Works in streaming setting. Faster & less memory intensive.
- Can use (almost) any “exact counting” method for step 3.
[Paranjape-Benson-Leskovec 17; Mackey+ 18; Liu-Benson-Charikar 18]
- Can parallelize over shifts and sampling windows
→ exposes parallelism to otherwise sequential algorithms.

THANKS!

Slides. `tinyurl.com/wsdm19`

Paper. *arXiv:1810.00980*

✉ `paul.liu@stanford.edu`



Cornell University



Stanford
University


Existing methods for temporal analysis are insufficient.

1. Models for network growth

Growth of academic collaborations, Internet infrastructure, etc. [Leskovec+ 07]

2. Sequence of snapshot aggregates

Daily phone call graph [Araujo+ 14], weekly email snapshots [Xu-Hero 14]



Modern temporal network datasets

- fine-scale time resolution
- high-frequency
- many repeated edges